

```
In [41]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from scipy import stats
from sklearn.preprocessing import LabelEncoder
from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score
from sklearn.metrics import classification_report
# Set a random seed for reproducibility
np.random.seed(42)
```

```
In [2]: # Generate synthetic data
num_customers = 1000
```

```
In [3]: # Generate customer IDs
customer_ids = np.arange(1, num_customers + 1)
```

```
In [4]: # Generate customer activities (time spent on different sections of the website)
activities = ['Home', 'Products', 'Blog', 'About', 'Contact']
num_activities = len(activities)
customer_activities = np.random.randint(low=1, high=60, size=(num_customers, num_activities))
```

```
In [5]: # Generate visit frequencies
visit_frequencies = np.random.randint(low=1, high=50, size=num_customers)
```

```
In [6]: # Generate advertisement categories
advertisement_categories = ['Sport-based', 'Cinematic', 'Artistic', 'Technology', 'Nature', 'Education']
num_categories = len(advertisement_categories)
customer_categories = np.random.choice(advertisement_categories, size=num_customers)
```

```
In [7]: # Create the dataset
data = pd.DataFrame({
    'CustomerID': customer_ids,
    'Home': customer_activities[:, 0],
    'Products': customer_activities[:, 1],
    'Blog': customer_activities[:, 2],
    'About': customer_activities[:, 3],
    'Contact': customer_activities[:, 4],
    'VisitFrequency': visit_frequencies,
    'Category': customer_categories
})
```

```
In [8]: # Save the dataset to a CSV file
data.to_csv('customer_data.csv', index=False)
```

```
In [9]: # Step 1: Importing the dataset
data = pd.read_csv('customer_data.csv')
```

```
In [10]: data
```

```
Out[10]:
```

	CustomerID	Home	Products	Blog	About	Contact	VisitFrequency	Category
0	1	39	52	29	15	43	4	Fashion
1	2	8	21	39	58	19	40	Technology
2	3	23	11	11	24	53	3	Cinematic
3	4	36	40	24	3	22	29	Sport-based
4	5	53	2	24	44	30	41	Artistic
...
995	996	7	52	57	12	16	6	Technology
996	997	16	43	53	36	13	21	Sport-based
997	998	7	36	47	51	9	32	Cinematic
998	999	32	11	8	33	23	35	Technology
999	1000	25	56	8	45	22	15	Sport-based

1000 rows × 8 columns

```
In [11]: # Step 2: Pre-processing the data
# Check for missing values
missing_values = data.isnull().sum()
print('Missing Values:')
print(missing_values)
```

```
Missing Values:
CustomerID      0
Home            0
Products        0
Blog            0
About           0
Contact         0
VisitFrequency  0
Category        0
dtype: int64
```

```
In [12]: # Check data types
data_types = data.dtypes
print('\nData Types:')
print(data_types)
```

Data Types:

```
CustomerID      int64
Home            int64
Products        int64
Blog            int64
About           int64
Contact         int64
VisitFrequency  int64
Category        object
dtype: object
```

```
In [13]: # Step 3: Exploratory Data Analysis (EDA)
# Summary statistics
summary_stats = data.describe()
print('\nSummary Statistics:')
print(summary_stats)
```

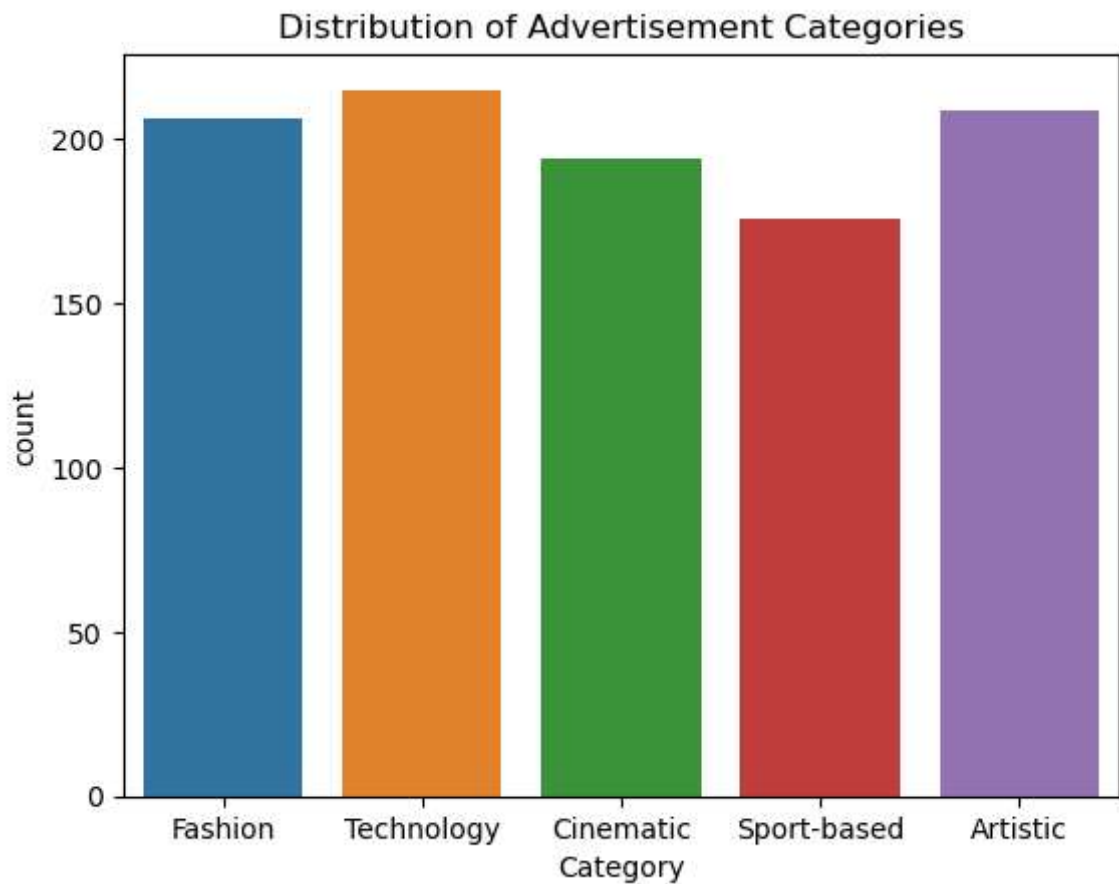
Summary Statistics:

	CustomerID	Home	Products	Blog	About \
count	1000.000000	1000.000000	1000.000000	1000.000000	1000.000000
mean	500.500000	30.532000	28.821000	29.972000	30.689000
std	288.819436	17.209915	16.770039	16.236096	17.187952
min	1.000000	1.000000	1.000000	1.000000	1.000000
25%	250.750000	16.000000	14.000000	17.000000	15.000000
50%	500.500000	31.000000	28.000000	30.000000	32.000000
75%	750.250000	46.000000	44.000000	43.250000	45.000000
max	1000.000000	59.000000	59.000000	59.000000	59.000000


	Contact	VisitFrequency
count	1000.000000	1000.000000
mean	30.615000	25.154000
std	17.088323	14.381346
min	1.000000	1.000000
25%	16.000000	13.000000
50%	31.000000	25.500000
75%	46.000000	38.000000
max	59.000000	49.000000

```
In [14]: # Distribution of the target variable
sns.countplot(data['Category'])
plt.title('Distribution of Advertisement Categories')
plt.show()
```

C:\Users\DELL\anaconda3\lib\site-packages\seaborn_decorators.py:36: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.
warnings.warn(



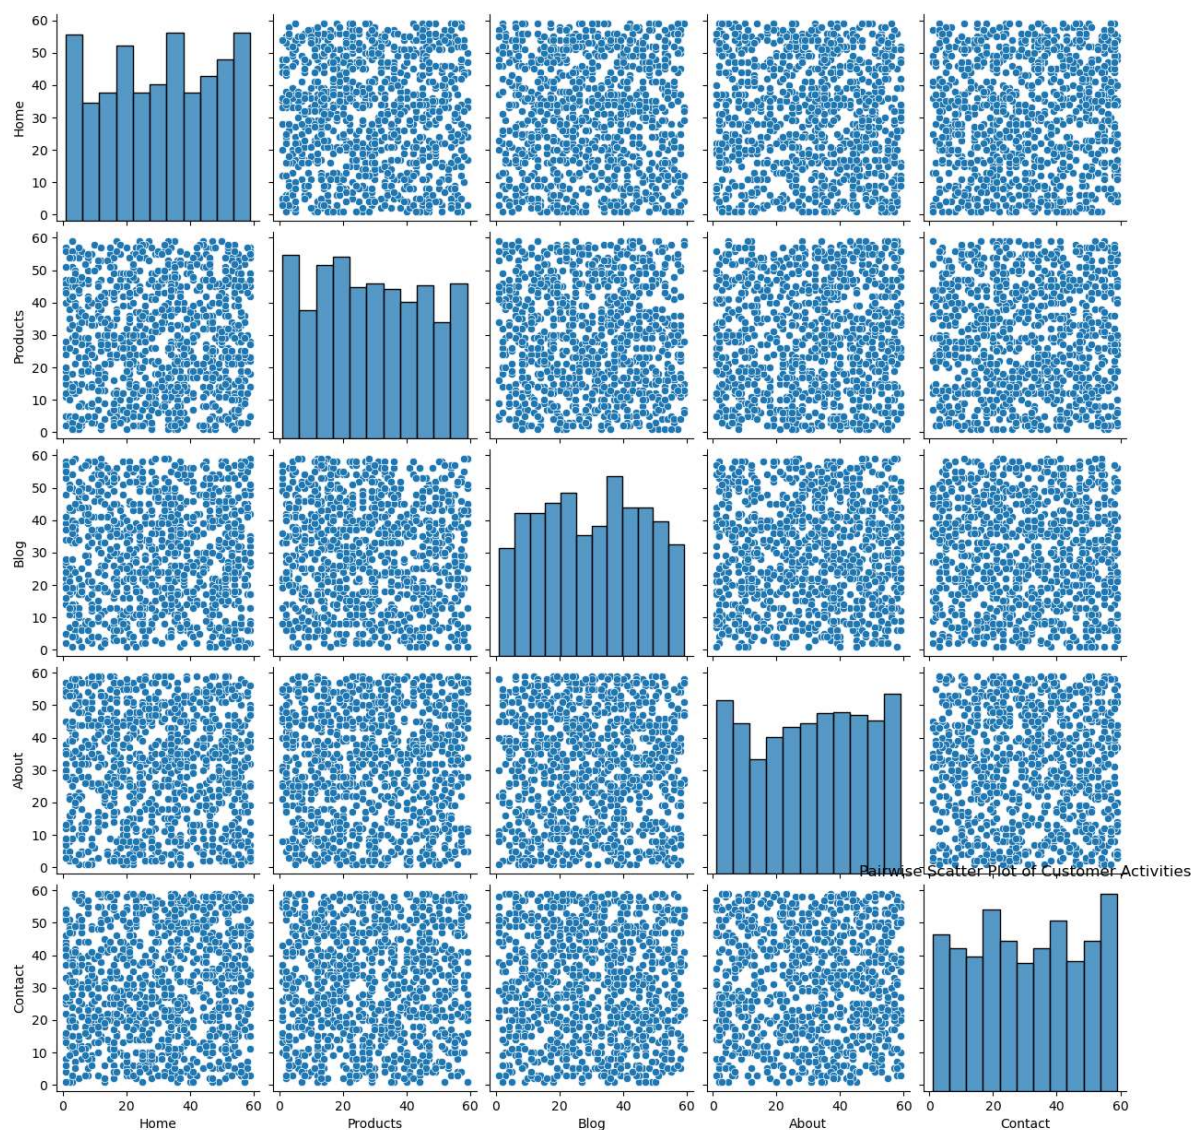
```
In [15]: # Correlation matrix
correlation_matrix = data.drop('CustomerID', axis=1).corr()
plt.figure(figsize=(10, 8))
sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm')
plt.title('Correlation Matrix')
plt.show()
```



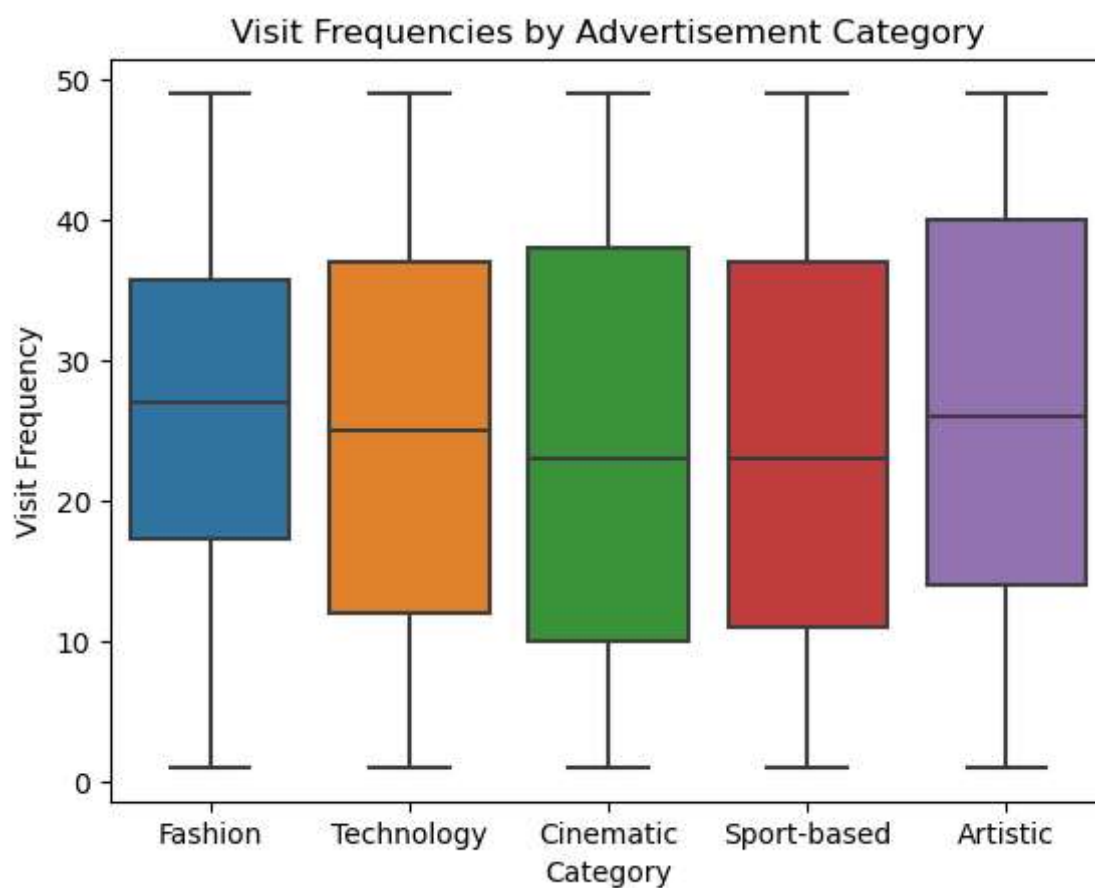
```
In [16]: # Step 4: Statistical Analysis
# Perform t-test for a specific feature (e.g., Home vs. Products)
feature1 = data[data['Category'] == 'Sport-based']['Home']
feature2 = data[data['Category'] == 'Sport-based']['Products']
t_stat, p_value = stats.ttest_ind(feature1, feature2)
print('\nT-test Results:')
print('t-statistic:', t_stat)
print('p-value:', p_value)
```

T-test Results:
t-statistic: 1.6481839294931617
p-value: 0.10021227176690972

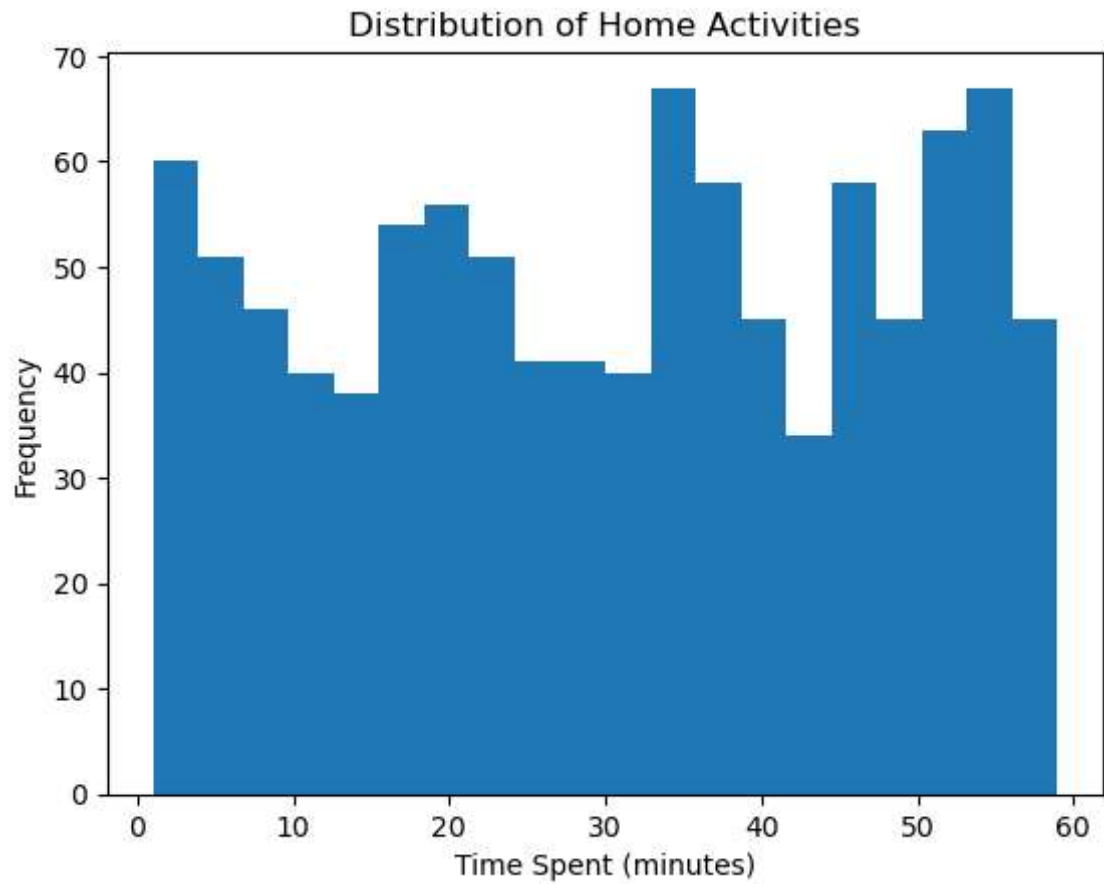
```
In [17]: # Pairwise Scatter Plot of Customer Activities
sns.countplot(data['Category'])
sns.pairplot(data[['Home', 'Products', 'Blog', 'About', 'Contact']])
plt.title('Pairwise Scatter Plot of Customer Activities')
plt.show()
```




```
In [18]: # Box Plot of Visit Frequencies by Advertisement Category
sns.boxplot(x='Category', y='VisitFrequency', data=data)
plt.title('Visit Frequencies by Advertisement Category')
plt.xlabel('Category')
plt.ylabel('Visit Frequency')
plt.show()
```



```
In [19]: # Histogram of Home Activities
plt.hist(data['Home'], bins=20)
plt.title('Distribution of Home Activities')
plt.xlabel('Time Spent (minutes)')
plt.ylabel('Frequency')
plt.show()
```




```
In [20]: import pandas as pd
from scipy.stats import f_oneway

# Step 1: Importing the dataset
data = pd.read_csv('customer_data.csv')

# Step 2: Statistical Analysis - ANOVA
# Perform ANOVA on the 'VisitFrequency' feature across different categories
categories = data['Category'].unique()
category_data = [data[data['Category'] == category]['VisitFrequency'] for cate

# Perform the ANOVA test
f_stat, p_value = f_oneway(*category_data)

# Print the ANOVA results
print('ANOVA Results:')
print('F-statistic:', f_stat)
print('p-value:', p_value)
```

ANOVA Results:
F-statistic: 1.2052808152614456
p-value: 0.3068864688238199

```
In [21]: # Step 3: Specifying the problem scope
X = data.drop(['CustomerID', 'Category'], axis=1)
y = data['Category']
```

```
In [25]: # Step 4: Label Encoding the categorical column
label_encoder = LabelEncoder()
y_encoded = label_encoder.fit_transform(y)
```

```
In [26]: # Step 5: Choosing the most appropriate algorithm
# Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y_encoded, test_size=0.1)
```

```
In [27]: # Initialize the Random Forest classifier
clf = RandomForestClassifier()
```

```
In [28]: # Fit the classifier to the training data
clf.fit(X_train, y_train)
```

Out[28]: RandomForestClassifier()

```
In [29]: # Step 6: Evaluating the model
# Make predictions on the test data
y_pred = clf.predict(X_test)
```

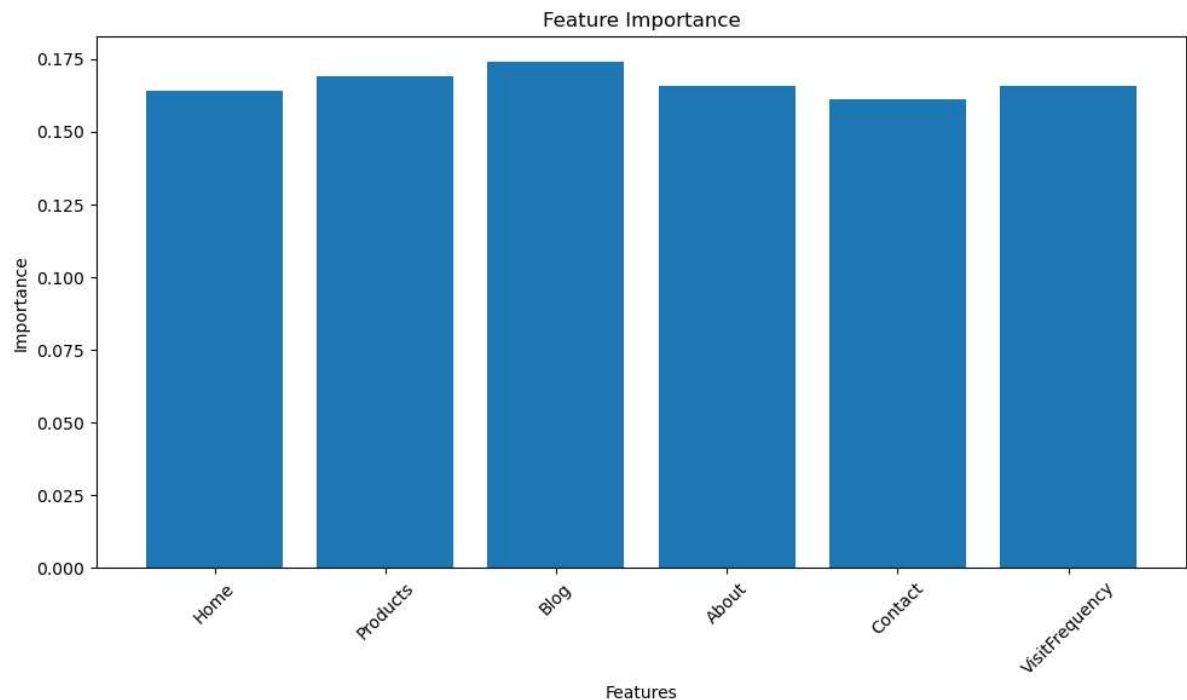
```
In [30]: # Calculate the evaluation metrics
accuracy = accuracy_score(y_test, y_pred)
precision = precision_score(y_test, y_pred, average='weighted')
recall = recall_score(y_test, y_pred, average='weighted')
f1 = f1_score(y_test, y_pred, average='weighted')

# Print the evaluation metrics
print('Evaluation Metrics:')
print('Accuracy:', accuracy)
print('Precision:', precision)
print('Recall:', recall)
print('F1-Score:', f1)
```

Evaluation Metrics:
Accuracy: 0.24
Precision: 0.24282102400523453
Recall: 0.24
F1-Score: 0.23895042197883243

```
In [32]: # Step 6: Calculating feature importance
feature_importance = clf.feature_importances_

# Step 7: Generating a bar plot for feature importance
plt.figure(figsize=(10, 6))
plt.bar(X.columns, feature_importance)
plt.xlabel('Features')
plt.ylabel('Importance')
plt.title('Feature Importance')
plt.xticks(rotation=45)
plt.tight_layout()
plt.show()
```



```
In [33]: # Define the model and hyperparameters to be tuned
model = RandomForestClassifier()
param_grid = {
    'n_estimators': [100, 200, 300],
    'max_depth': [None, 5, 10],
    'min_samples_split': [2, 5, 10]
}
```

```
In [36]: # Perform grid search to find the best hyperparameters
grid_search = GridSearchCV(model, param_grid, cv=5)
grid_search.fit(X_train, y_train)

# Retrieve the best model with tuned hyperparameters
best_model = grid_search.best_estimator_
```

```
In [37]: # Retrieve the best model with tuned hyperparameters
best_model = grid_search.best_estimator_
```

```
In [38]: # Step 7: Evaluating the model
# Make predictions on the test data
y_pred = best_model.predict(X_test)
```

```
In [39]: # Calculate the evaluation metrics
accuracy = accuracy_score(y_test, y_pred)
precision = precision_score(y_test, y_pred, average='weighted')
recall = recall_score(y_test, y_pred, average='weighted')
f1 = f1_score(y_test, y_pred, average='weighted')

# Print the evaluation metrics
print('Evaluation Metrics:')
print('Accuracy:', accuracy)
print('Precision:', precision)
print('Recall:', recall)
print('F1-Score:', f1)
```

```
Evaluation Metrics:
Accuracy: 0.26
Precision: 0.25936171552279225
Recall: 0.26
F1-Score: 0.2559636893250634
```

```
In [42]: # Step 8: Calculate evaluation metrics
report = classification_report(y_test, y_pred)

# Print the evaluation metrics
print('Evaluation Metrics:')
print(report)
```


Evaluation Metrics:

	precision	recall	f1-score	support
0	0.24	0.33	0.28	42
1	0.22	0.16	0.18	38
2	0.30	0.34	0.32	41
3	0.29	0.24	0.26	29
4	0.26	0.22	0.24	50
accuracy			0.26	200
macro avg	0.26	0.26	0.26	200
weighted avg	0.26	0.26	0.26	200

```
In [51]: from sklearn.metrics import confusion_matrix

cm = confusion_matrix(y_test, y_pred)
sns.heatmap(cm, annot=True, fmt='d').set_title('Random Forest confusion matrix')
```

Out[51]: Text(0.5, 1.0, 'Random Forest confusion matrix')



```
In [43]: from sklearn.preprocessing import StandardScaler
from sklearn.cluster import KMeans
```

```
In [44]: # Step 1: Importing the dataset
data = pd.read_csv('customer_data.csv')
# Step 2: Specifying the problem scope
X = data.drop(['CustomerID', 'Category'], axis=1)
```

```
In [45]: # Step 3: Standardize the features
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)
```

```
In [46]: # Step 4: Apply the K-means clustering algorithm
kmeans = KMeans(n_clusters=5, random_state=42)
clusters = kmeans.fit_predict(X_scaled)
```

```
In [47]: # Step 6: Analyze the clusters
cluster_df = pd.DataFrame({'CustomerID': data['CustomerID'], 'Cluster': clusters})

# Step 7: Print the cluster assignments
print(cluster_df)
```

	CustomerID	Cluster
0	1	0
1	2	4
2	3	1
3	4	2
4	5	3
..
995	996	1
996	997	4
997	998	4
998	999	2
999	1000	4

[1000 rows x 2 columns]

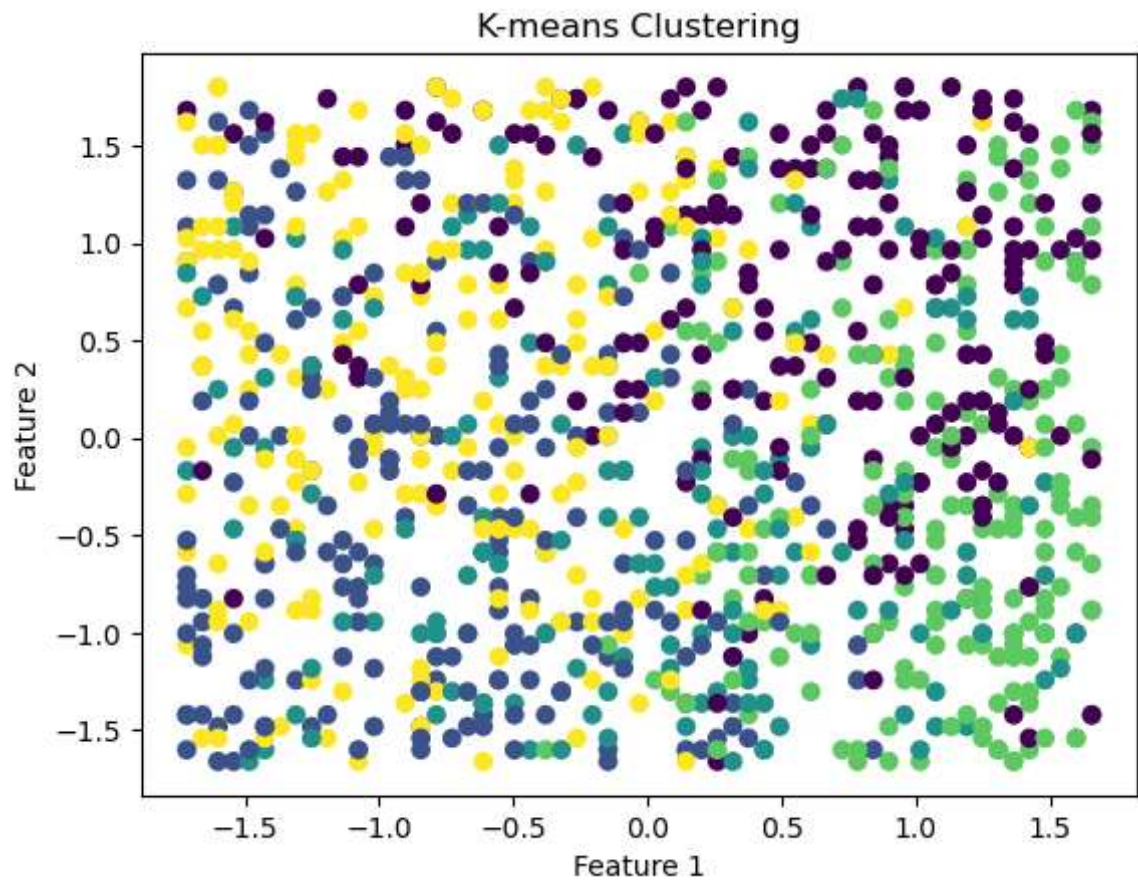
```
In [48]: from sklearn.metrics import silhouette_score, calinski_harabasz_score, davies_bouldin_score
# Step 7: Calculate evaluation metrics
silhouette = silhouette_score(X_scaled, clusters)
calinski_harabasz = calinski_harabasz_score(X_scaled, clusters)
davies_bouldin = davies_bouldin_score(X_scaled, clusters)

# Step 8: Print the evaluation metrics
print('Evaluation Metrics:')
print('Silhouette Score:', silhouette)
print('Calinski-Harabasz Index:', calinski_harabasz)
print('Davies-Bouldin Index:', davies_bouldin)
```

Evaluation Metrics:
 Silhouette Score: 0.1247273715235876
 Calinski-Harabasz Index: 119.84835524566272
 Davies-Bouldin Index: 1.8319984453825473

```
In [49]: import matplotlib.pyplot as plt

# Step 7: Visualize the clustering results
plt.scatter(X_scaled[:, 0], X_scaled[:, 1], c=clusters, cmap='viridis')
plt.xlabel('Feature 1')
plt.ylabel('Feature 2')
plt.title('K-means Clustering')
plt.show()
```



```
In [53]: from sklearn.svm import SVC
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score
from sklearn.model_selection import train_test_split

# Step 2: Create an SVM classifier
svm_classifier = SVC(kernel='linear')

# Step 3: Train the classifier on the training data
svm_classifier.fit(X_train, y_train)

# Step 4: Make predictions on the testing data
y_pred = svm_classifier.predict(X_test)

# Step 5: Evaluate the performance of the classifier
accuracy = accuracy_score(y_test, y_pred)
precision = precision_score(y_test, y_pred, average='weighted')
recall = recall_score(y_test, y_pred, average='weighted')
f1 = f1_score(y_test, y_pred, average='weighted')

# Step 6: Print the evaluation metrics
print('Evaluation Metrics:')
print('Accuracy:', accuracy)
print('Precision:', precision)
print('Recall:', recall)
print('F1-Score:', f1)
```

Evaluation Metrics:

Accuracy: 0.225

Precision: 0.2308424908424908

Recall: 0.225

F1-Score: 0.21889383896154535


```
In [54]: from sklearn.ensemble import RandomForestRegressor
from sklearn.metrics import mean_squared_error, mean_absolute_error
from sklearn.model_selection import train_test_split

# Step 2: Create a Random Forest Regression model
rf_regressor = RandomForestRegressor(n_estimators=100, random_state=42)

# Step 3: Train the model on the training data
rf_regressor.fit(X_train, y_train)

# Step 4: Make predictions on the testing data
y_pred = rf_regressor.predict(X_test)

# Step 5: Evaluate the performance of the model
mse = mean_squared_error(y_test, y_pred)
mae = mean_absolute_error(y_test, y_pred)

# Step 6: Print the evaluation metrics
print('Evaluation Metrics:')
print('Mean Squared Error:', mse)
print('Mean Absolute Error:', mae)
```

```
Evaluation Metrics:
Mean Squared Error: 2.3292535
Mean Absolute Error: 1.3198500000000002
```

In []: